



# Figure 1: workload sensitivity maps

Paper plot




Artifact result




# Figure 3: 128 MiB AllReduce

## Paper plot




## Artifact result




Maximum relative errors: 0.078% and 0.166% for latency; 2.373% and 1.493% for bandwidth.

# Figure 4: mixed collectives

Paper plot




Artifact result




# Figure 5: Llama training iteration

Paper plot




Artifact result




Maximum relative error 0.014853%; mean relative error 0.004879% over 201 points.

# Figure 6: latency and bandwidth sensitivity

## Paper plot



## Artifact result



**Figure 7: monolithic memory footprint**

Paper plot






Figure 7: Monolithic LP's peak memory footprint for  
Composition LP's peak memory footprint for

Artifact result



# Figure 8: network-parameter sensitivity

Paper plot



Artifact result






Figure 9: network cost trajectory

Paper plot





Artifact result





# Figure 10: jitter sensitivity

Paper plot



Artifact result

