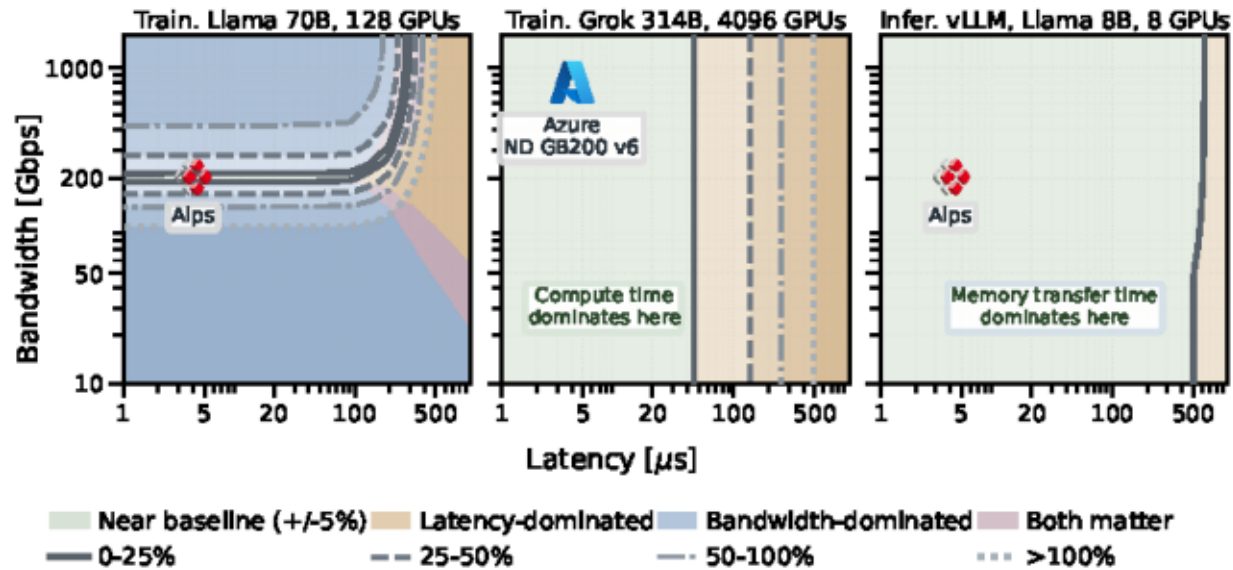


Figure 1: workload sensitivity maps

Paper plot



Artifact result

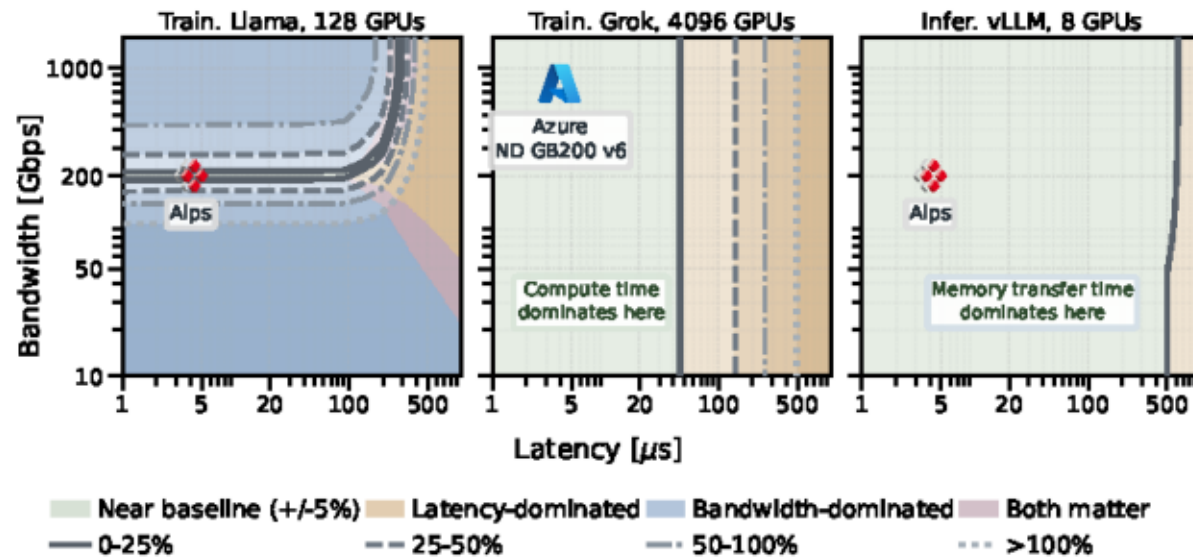
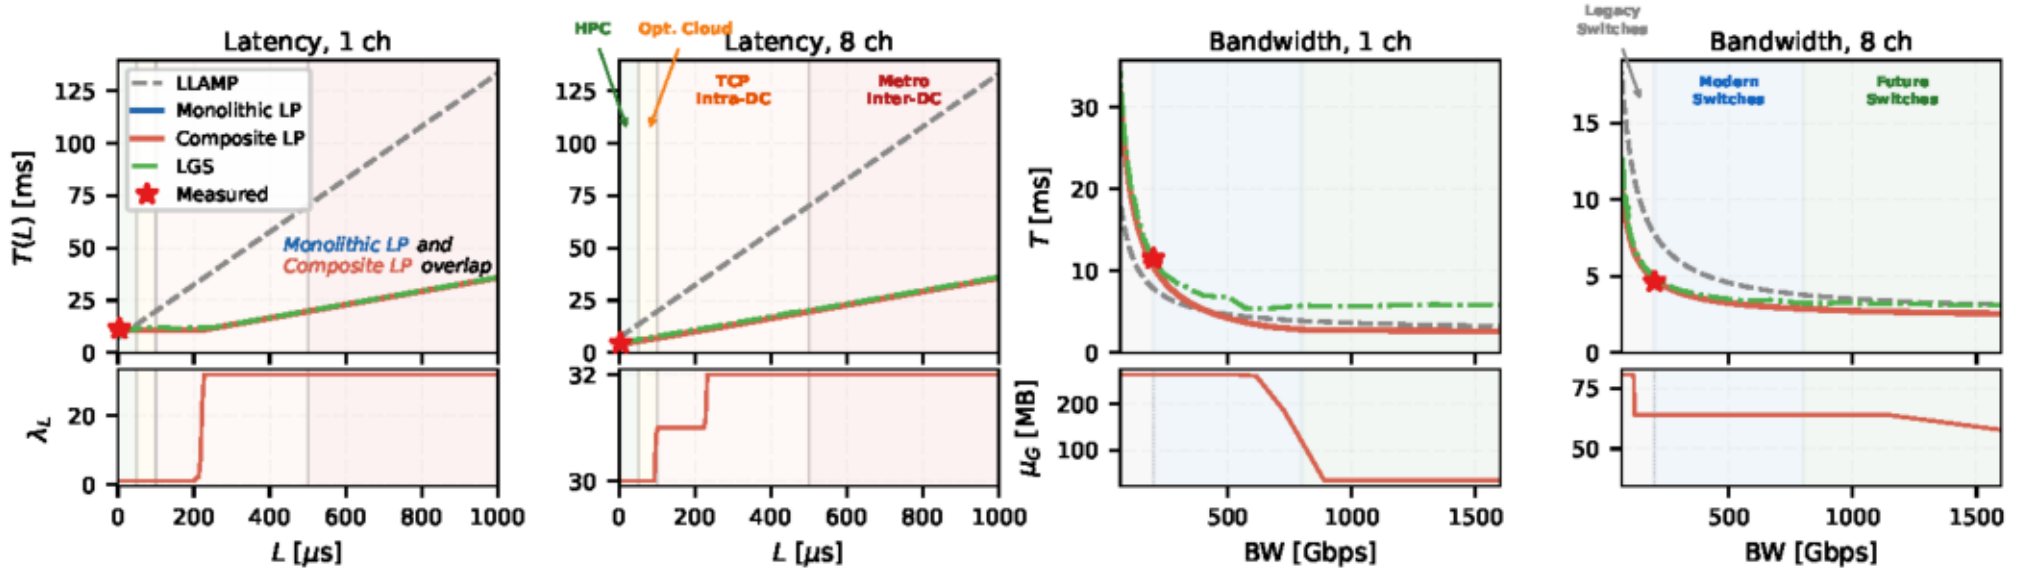


Figure 3: 128 MiB AllReduce

Paper plot



Artifact result

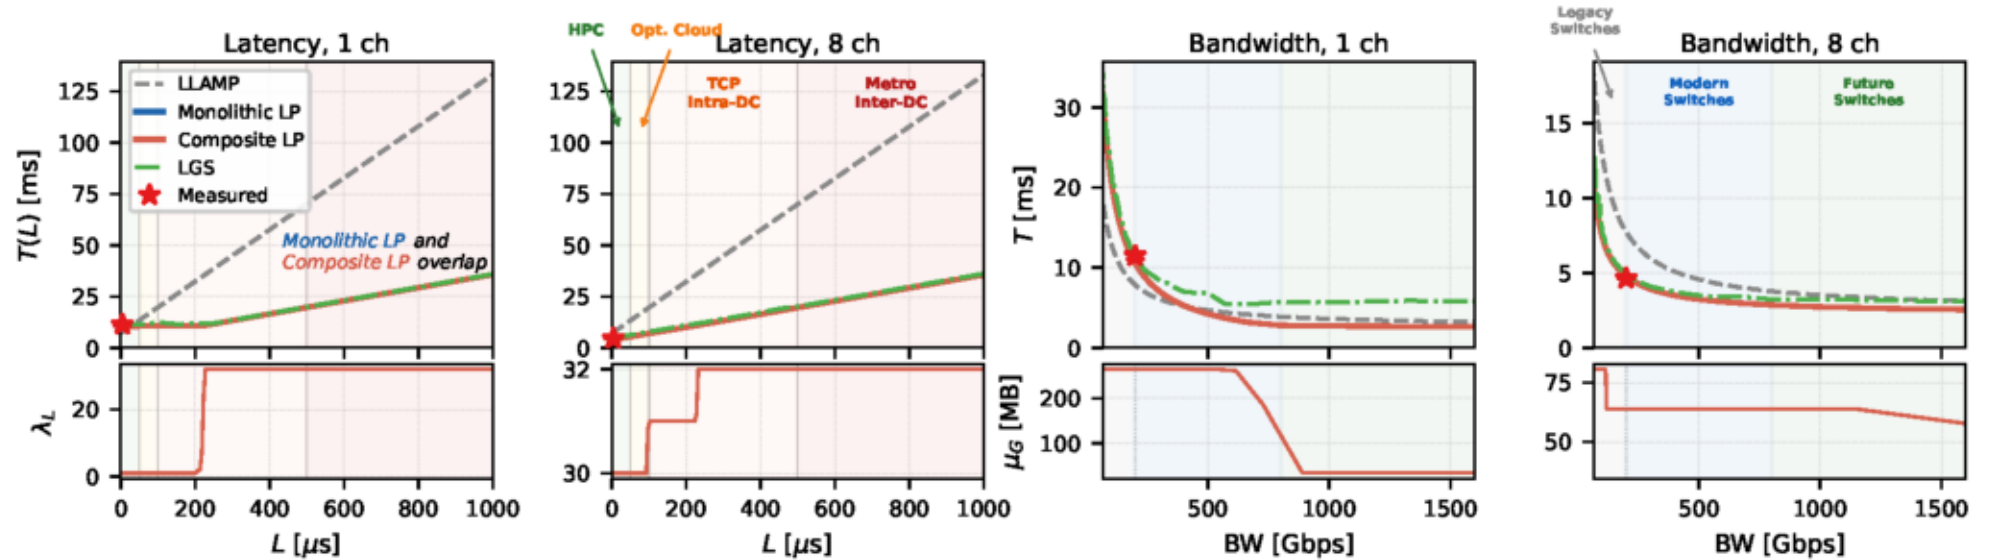
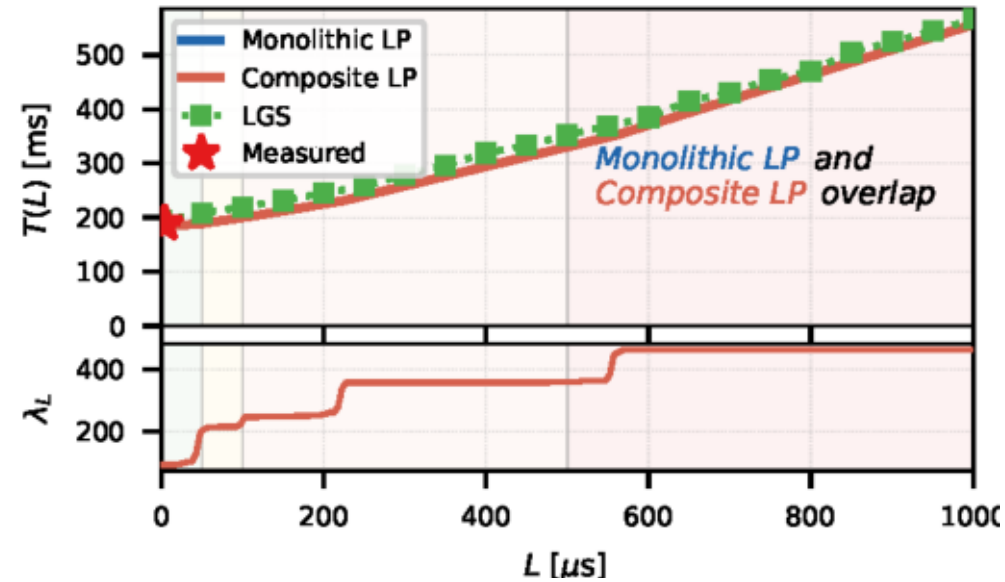


Figure 4: mixed collectives

Paper plot



Artifact result

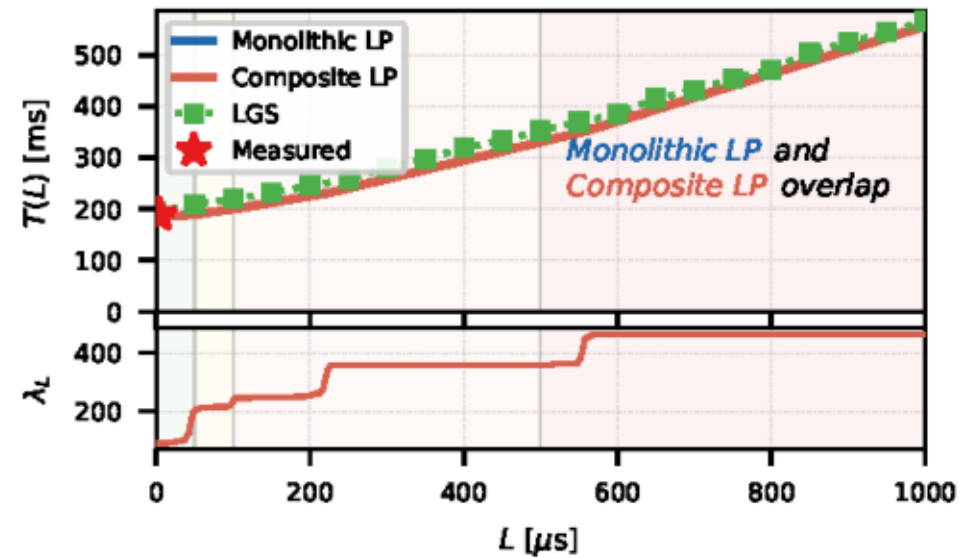
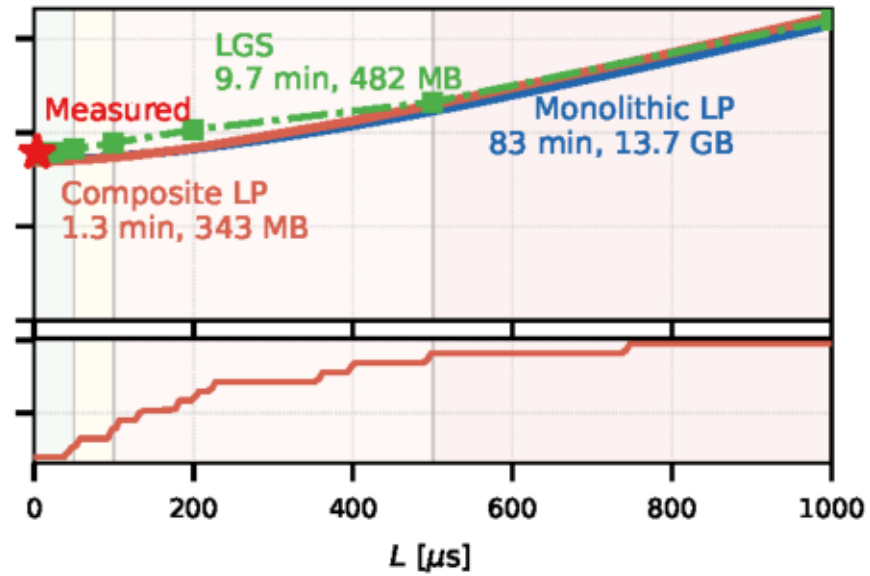


Figure 5: Llama training iteration

Paper plot



Artifact result

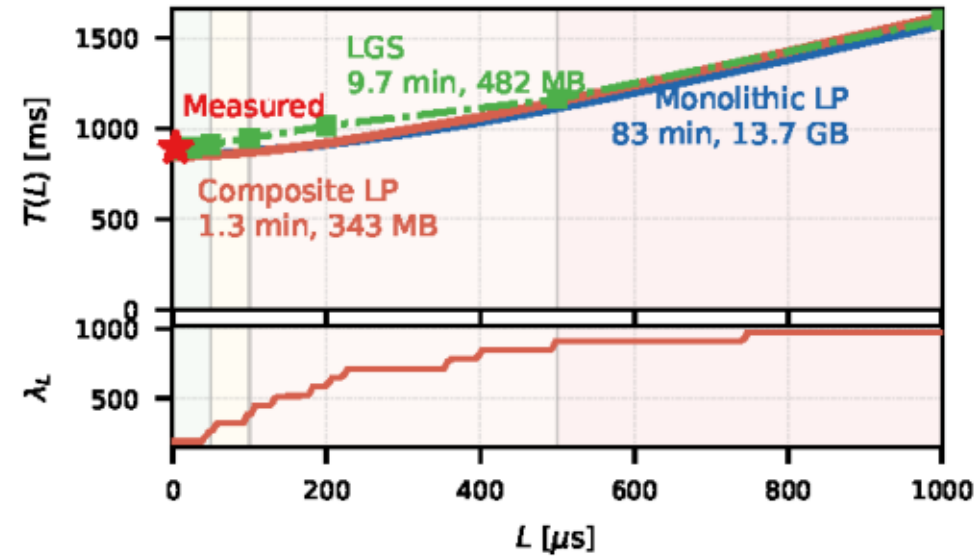
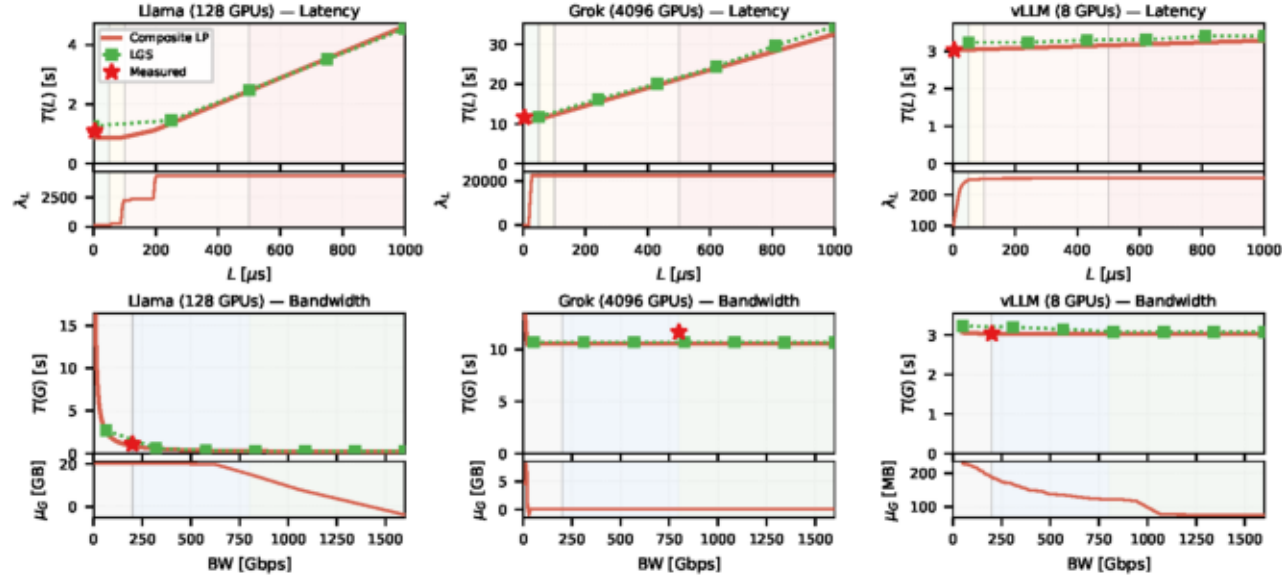


Figure 6: latency and bandwidth sensitivity

Paper plot



Artifact result

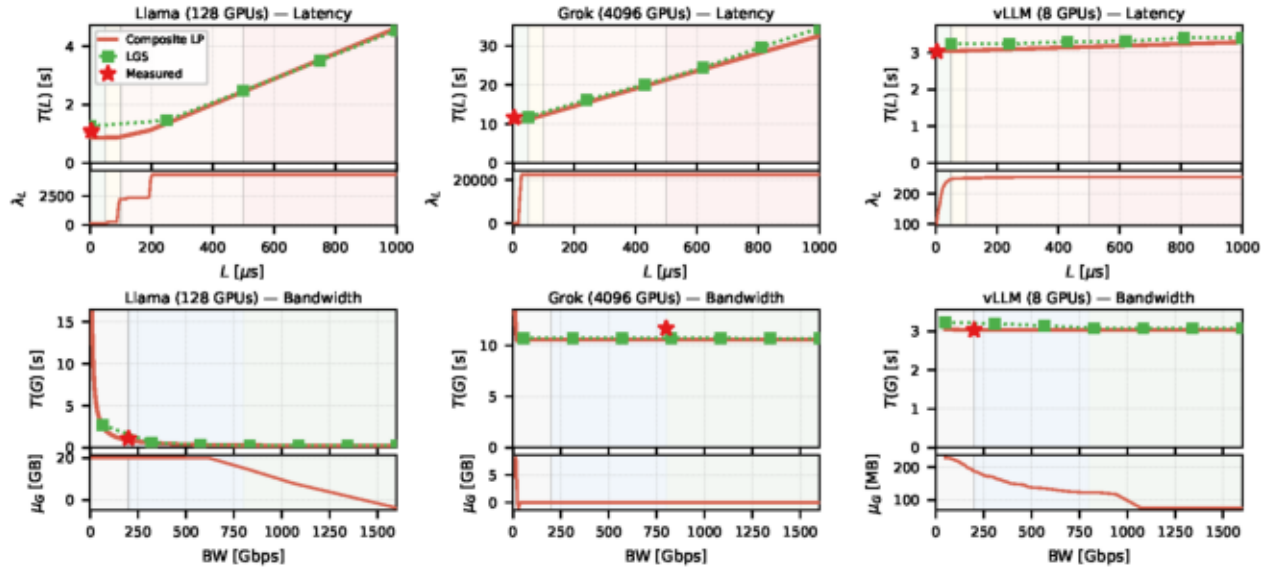


Figure 7: monolithic memory footprint

Paper plot

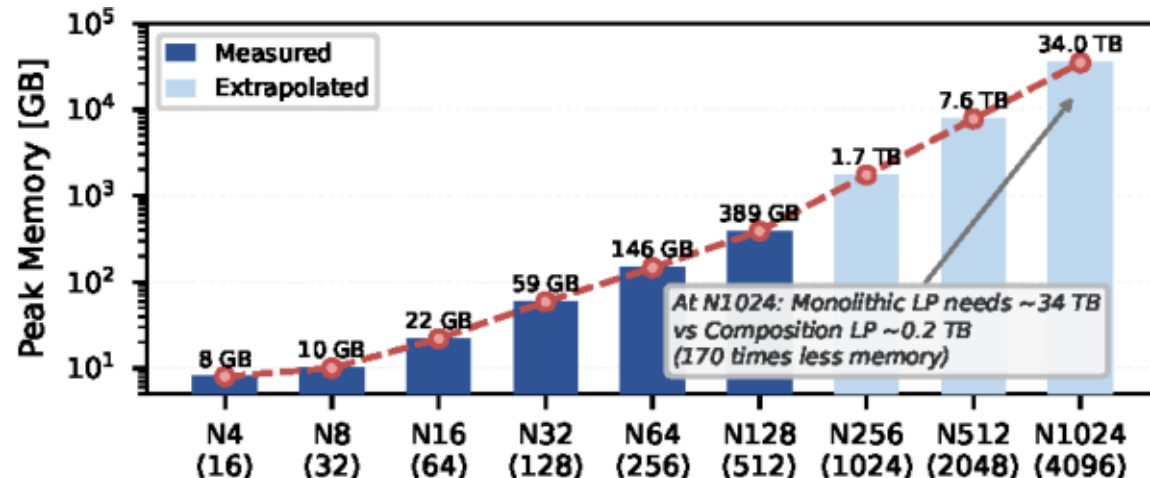


Figure 7: Monolithic LP's peak memory footprint for
Compositional LP's peak memory footprint for

Artifact result

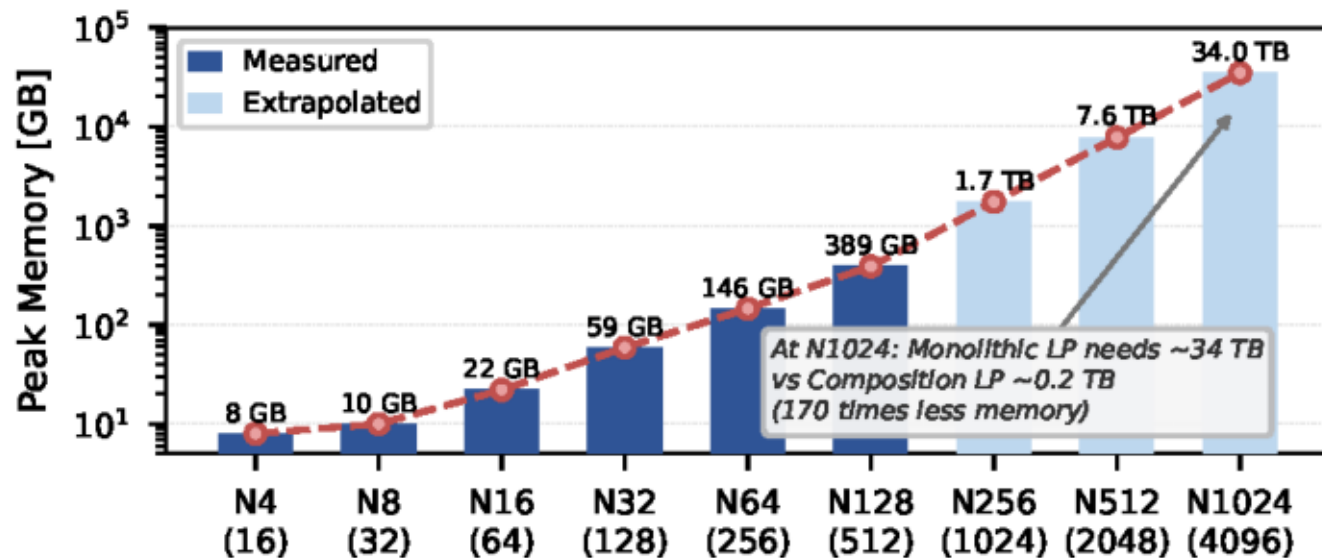
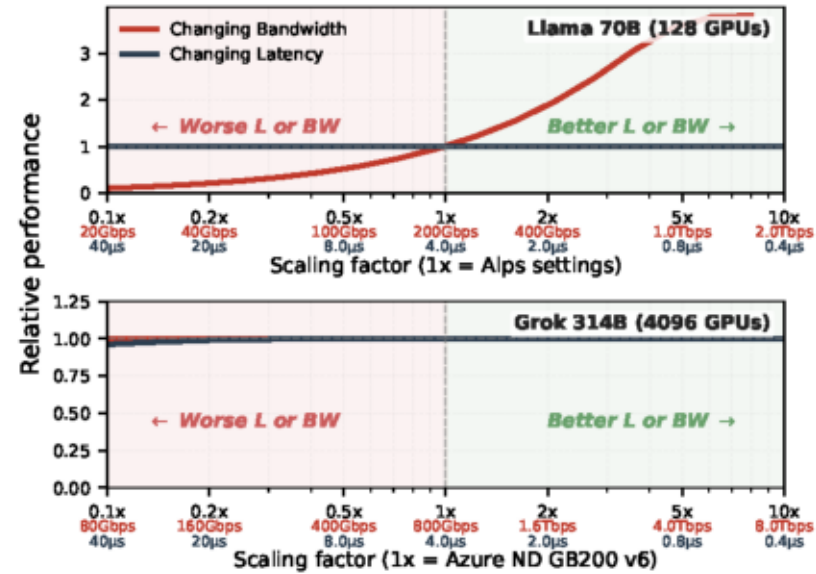


Figure 8: network-parameter sensitivity

Paper plot



Artifact result

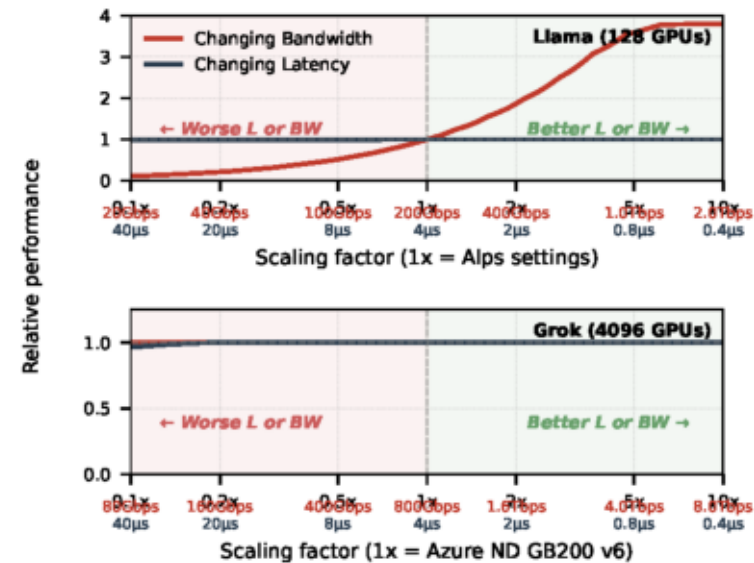
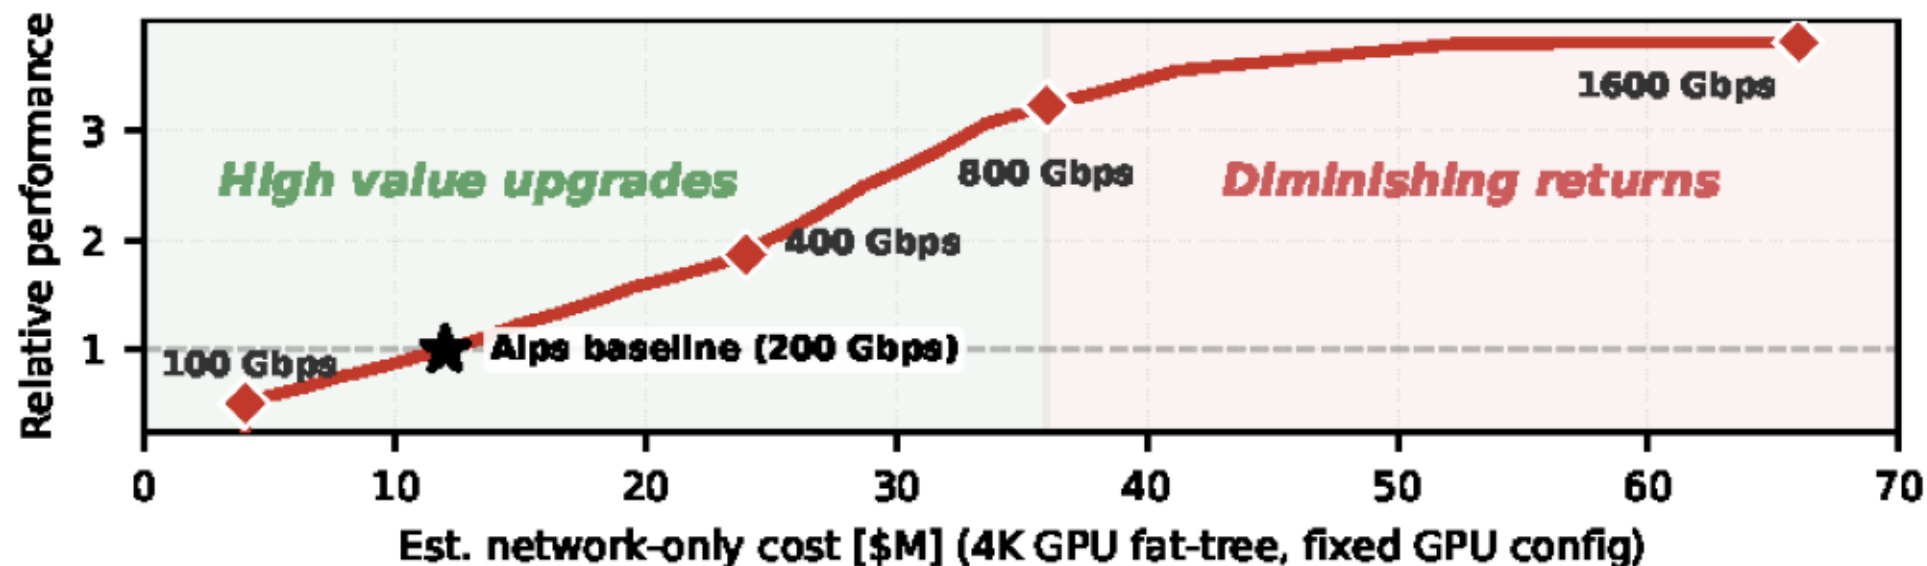


Figure 9: network cost trajectory

Paper plot



Artifact result

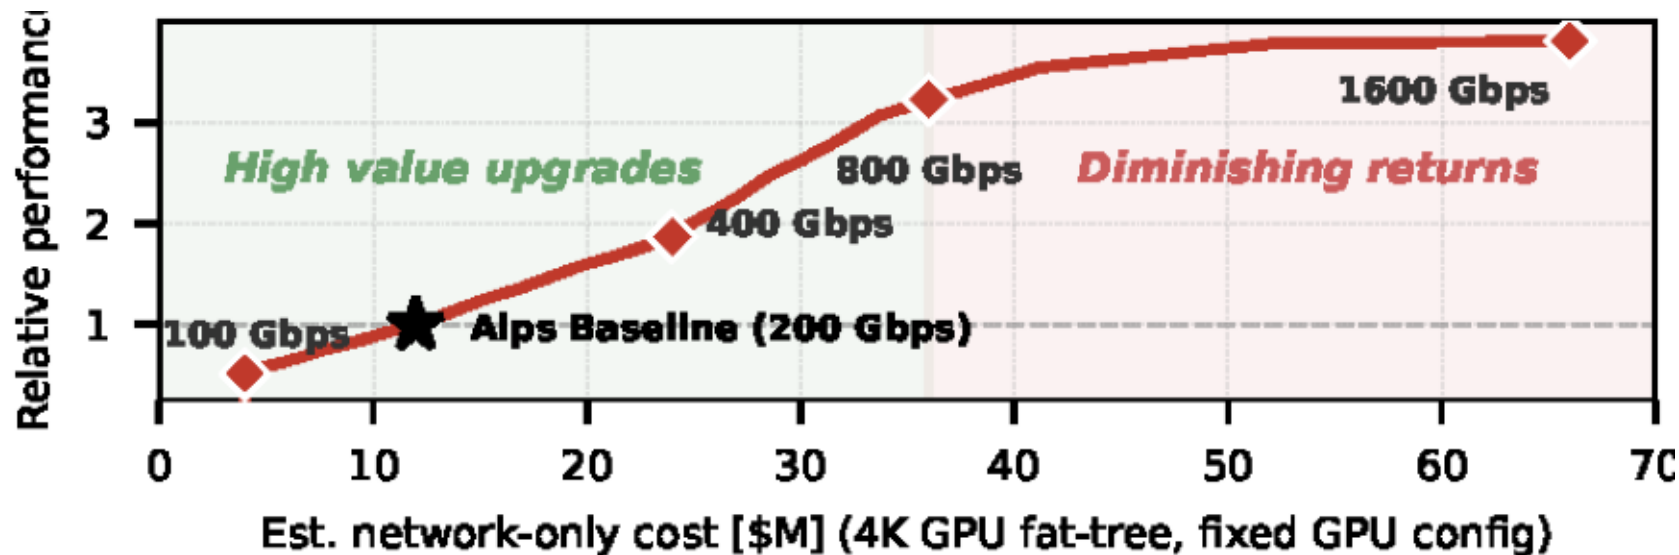
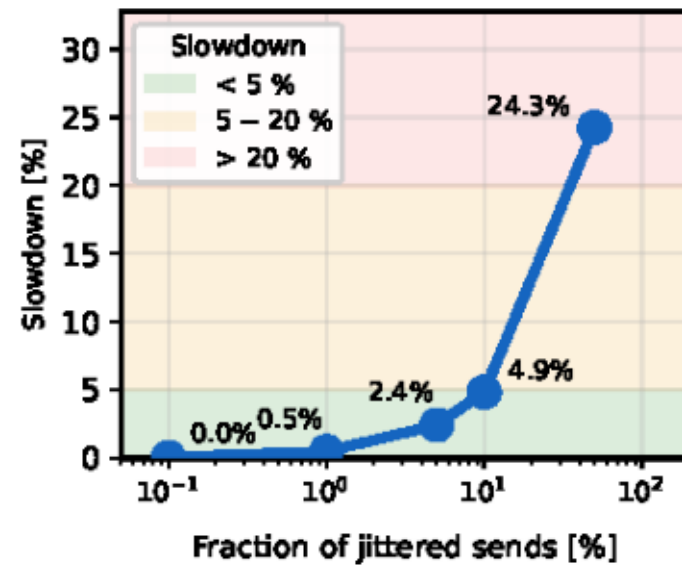
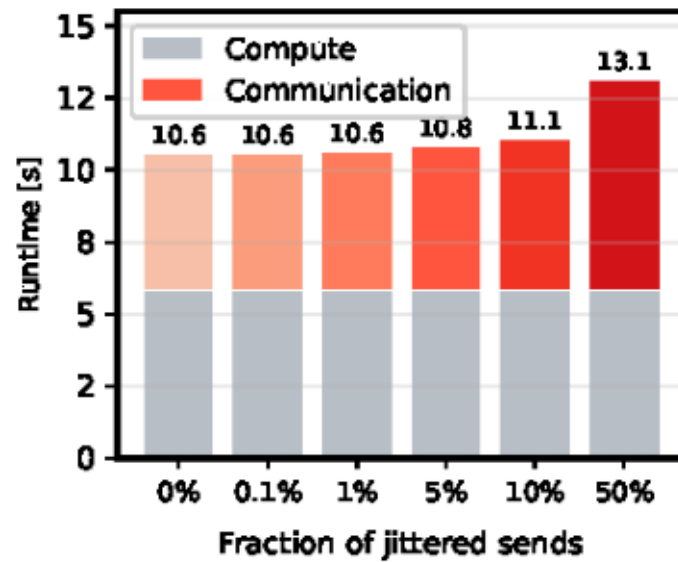


Figure 10: jitter sensitivity

Paper plot



Artifact result

