




Llama (128 GPUs) — Latency




Grok (4096 GPUs) — Latency




vLLM (8 GPUs) — Latency



Llama (128 GPUs) — Bandwidth



Grok (4096 GPUs) — Bandwidth



vLLM (8 GPUs) — Bandwidth

